



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Lecturer
Justin Yokota

C Pointers, Arrays, and Strings

Computing in the News

Hackers Can Make Computers Destroy Their Own Chips with Electricity (New Scientist, 2023-01-19)

<https://www.newscientist.com/article/2354844-hackers-can-make-computers-destroy-their-own-chips-with-electricity/>



Zitai Chen and David Oswald at the U.K.'s University of Birmingham uncovered a bug in the control systems of server motherboards that could be exploited to compromise sensitive information or to **destroy their central processing units (CPUs)**. The researchers found a feature in the Supermicro X11SSL-CF motherboard often used in servers that they could tap to upload their own control software. Chen and Oswald discovered a flash memory chip in the motherboard's baseboard management controller that **they could remotely command to send excessive electrical current through the CPU, destroying it in seconds**. After the researchers disclosed the flaw to Supermicro, the company said it has rated its severity as "high" and has patched the bug in its existing motherboards.

Address vs. Value

- **Consider memory to be a single huge array:**
 - Each cell of the array has an address associated with it.
 - Each cell also stores some value.
 - Do you think they use signed or unsigned numbers?
Negative address?!
- **Don't confuse the **address** referring to a memory location with the **value** stored in that location.**
- **For now, the abstraction lets us think we have access to ∞ memory, numbered from 0...**



Pointers

- An address refers to a particular memory location. In other words, it points to a memory location.
- **Pointer**: A variable that contains the address of a variable.



Pointer Syntax

- `int *p;`
 - Tells compiler that **variable `p`** is **address** of an `int`


- `p = &y;`
 - Tells compiler to assign **address of `y`** to `p`
 - `&` called the “address operator” in this context

- `z = *p;`
 - Tells compiler to assign **value at address in `p`** to `z`
 - `*` called the “dereference operator” in this context

Pointers

- How to create a pointer:

& operator: get address of a variable



Note the “*” gets used 2 different ways in this example. In the declaration to indicate that **p** is going to be a pointer, and in the **printf** to get the value pointed to by **p**.

- How get a value pointed to?

* “dereference operator”: get value pointed to

```
printf("p points to %d\n", *p);
```

- **How to change a variable pointed to?**
 - Use dereference `*` operator on left of `=`



Pointers and Parameter Passing (1/2)

- **Java and C pass parameters “by value”**
 - procedure/function/method gets a copy of the parameter, so changing the copy cannot change the original

```
void addOne (int x) {
    x = x + 1;
}

int y = 3;
addOne (y) ;
```

y is still = 3

Pointers and Parameter Passing (2/2)

- How to get a function to change a value?

```
void addOne (int *p) {
    *p = *p + 1;
}

int y = 3;
addOne (&y) ;
```

y is now = 4

More C Pointer Dangers

- Declaring a pointer just allocates space to hold the pointer – it does not allocate something to be pointed to!
- **Local variables in C are not initialized**, they may contain anything.
- What does the following code do?

```
void f()  
{  
    int *ptr;  
    *ptr = 5;  
}
```

Pointers in C ... The Good, Bad, and the Ugly

- **Why use pointers?**
 - If we want to pass a large struct or array, it's easier / faster / etc. to pass a pointer than the whole thing
 - Otherwise we'd need to copy a huge amount of data
 - In general, pointers allow cleaner, more compact code
- **So what are the drawbacks?**
 - Pointers are probably the single largest source of bugs in C, so be careful anytime you deal with them
 - Most problematic with dynamic memory management—coming up next time
 - Dangling references and memory leaks



Using Pointers Effectively

- Pointers are used to point to **any** data type (`int`, `char`, a struct, etc.).
- Normally a pointer can only point to one type (`int`, `char`, a struct, etc.).
 - `void *` is a type that can point to anything (generic pointer)
 - Use sparingly to help avoid program bugs... and security issues... and a lot of other bad things!
- You can even have pointers to functions...
 - `int (*fn) (void *, void *) = &foo`
 - `fn` is a function that accepts two `void *` pointers and returns an `int` and is initially pointing to the function `foo`.
 - `(*fn) (x, y)` will then call the function

Pointers and Structures

```
typedef struct {
    int x;
    int y;
} Point;
```

```
Point p1;
Point p2;
Point *paddr;
```

```
/* dot notation */
int h = p1.x;
p2.y = p1.y;
```

```
/* arrow notation */
int h = paddr->x;
int h = (*paddr).x;
```

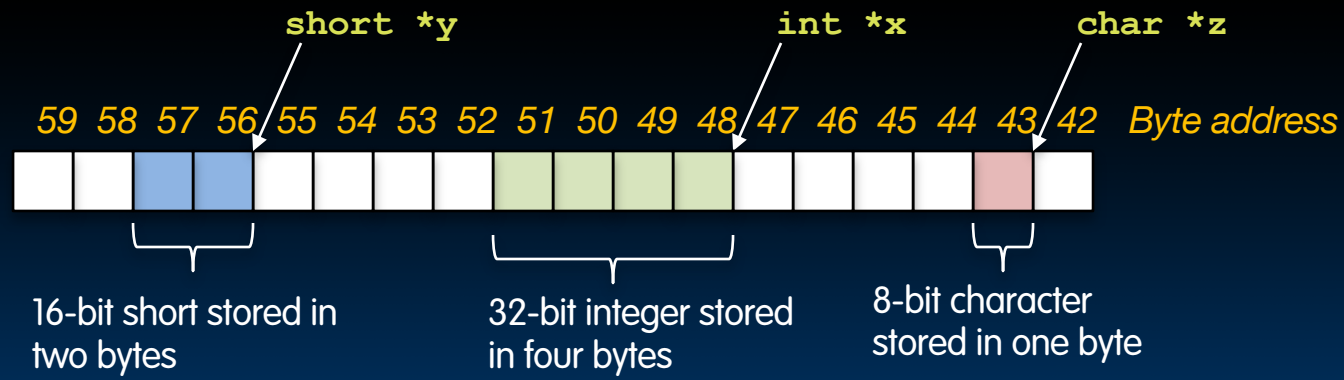
```
/* This works too */
p1 = p2;
```

NULL pointers...

- **The pointer of all 0s is special**
 - The "NULL" pointer, like in Java, python, etc...
- **If you write to or read a null pointer, your program should crash**
- **Since "0 is false", its very easy to do tests for null:**
 - `if(!p) { /* P is a null pointer */ }`
 - `if(q) { /* Q is not a null pointer */ }`

Pointing to Different Size Objects

- **Modern machines are “byte-addressable”**
 - Hardware’s memory composed of 8-bit storage cells, each has a unique address
- **A C pointer is just abstracted memory address**
- **Type declaration tells compiler how many bytes to fetch on each access through ptr**
 - E.g., 32-bit integer stored in 4 consecutive 8-bit bytes
- **But we actually want “word alignment”**
 - Some processors will not allow you to address 32b values without being on 4-byte boundaries
 - Others will just be very slow if you try to access “unaligned” memory.



Arrays

Arrays (1/5)

- **Declaration:**
 - `int ar[2];`
 - ...declares a 2-element integer array
 - An array is really just a block of memory
- **Declaration and initialization**
 - `int ar[] = {795, 635};`
 - declares and fills a 2-elt integer array
- **Accessing elements:**
 - `ar[num]`
 - returns the numth element.

Arrays (2/5)

- Arrays are (almost) identical to pointers
 - `char *string` and `char string[]` are nearly identical declarations
 - They differ in very subtle ways: incrementing, declaration of filled arrays
- **Key Concept:** An array variable is a “pointer” to the first element.

Arrays (3/5)

- **Consequences:**
 - **ar** is an array variable but looks like a pointer in many respects (though not all)
 - **ar[0]** is the same as ***ar**
 - **ar[2]** is the same as ***(ar+2)**
 - We can use pointer arithmetic to access arrays more conveniently.
- **Declared arrays are only allocated while the scope is valid**

```
char *foo() {
    char string[32]; ...;
    return string;
} is incorrect
```

Arrays (4/5)

- Array size n ; want to access from 0 to $n-1$, so you should use counter AND utilize a variable for declaration & incr

- Wrong

```
int i, ar[10];  
for(i = 0; i < 10; i++){ ... }
```

- Right

```
int ARRAY_SIZE = 10;  
int i, a[ARRAY_SIZE];  
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```

- Why? SINGLE SOURCE OF TRUTH
 - You're utilizing **indirection** and avoiding maintaining two copies of the number 10

Arrays (5/5)

- **Pitfall: An array in C does not know its own length, & bounds not checked!**
 - Consequence: We can accidentally access off the end of an array.
 - Consequence: We must pass the array and its size to a procedure which is going to traverse it.
- **Segmentation faults and bus errors:**
 - These are VERY difficult to find; be careful!
 - You'll learn how to debug these in lab...

Pointer Arithmetic

- **pointer + n**
 - Adds **$n * \text{sizeof}$** ("whatever pointer is pointing to") to the memory address
- **pointer - n**
 - Subtracts **$n * \text{sizeof}$** ("whatever pointer is pointing to") from the memory address

Pointers (1/4) ...review...

- **Java and C pass parameters “by value”**
 - procedure/function/method gets a copy of the parameter, so changing the copy cannot change the original

```
void addOne (int x) {  
    x = x + 1;  
}
```

```
int y = 3;  
addOne (y) ;
```

y is still = 3

Pointers (2/4) ...review...

- How to get a function to change a value?

```
void addOne (int *p) {  
    *p = *p + 1;  
}  
  
int y = 3;  
addOne (&y) ;
```


y is now = 4

Pointers (3/4)

- But what if you want to change a pointer?
 - What gets printed?

```
void IncrementPtr(int *p)
{    p = p + 1; }
```

```
int A[3] = {50, 60, 70};
int *q = A;
IncrementPtr( q );
printf( "*q = %d\n", *q );
```




Pointers (4/4)

- Idea! Pass **a pointer to a pointer!**
 - Declared as ****h**
 - Now what gets printed?

```
void IncrementPtr(int **h)
{    *h = *h + 1;    }
```

```
int A[3] = {50, 60, 70};
int *q = A;
IncrementPtr(&q);
printf("*q = %d\n", *q);
```



Clicker

L04 How many syntax+logic errors in this C18 code?

```
void main(); {  
    int *p, x=5, y; // init  
    y = *(p = &x) + 1;  
    int z;  
    flip-sign(p);  
    printf("x=%d,y=%d,p=%d\n",x,y,p);  
}  
flip-sign(int *n){*n = -(*n)}
```

0 A
1 B
2 C
3 D
4 E
5 F
6 G

Function Pointer Example



map (actually mutate_map easier)

```
#include <stdio.h>
```

```
int x10(int), x2(int);  
void mutate_map(int [], int n, int(*) (int));  
void print_array(int [], int n);
```

```
int x2 (int n) { return 2*n; }  
int x10(int n) { return 10*n; }
```

```
void mutate_map(int A[], int n, int(*)fp) (int)) {  
    for (int i = 0; i < n; i++)  
        A[i] = (*fp) (A[i]);  
}
```

```
void print_array(int A[], int n) {  
    for (int i = 0; i < n; i++)  
        printf("%d ", A[i]);  
    printf("\n");  
}
```

```
% ./map  
3 1 4  
6 2 8  
60 20 80
```

```
int main(void)  
{  
    int A[] = {3,1,4}, n = 3;  
    print_array(A, n);  
    mutate_map (A, n, &x2);  
    print_array(A, n);  
    mutate_map (A, n, &x10);  
    print_array(A, n);  
}
```

And In Conclusion, ...

- All data is in memory. Each memory location has an address to use to refer to it and a value stored in it.
- A pointer is a C version of the address.
 - * “follows” a pointer to its value
 - & gets the address of a value
 - Use handles to change pointers
- Arrays don’t know their own size and bounds are not checked, careful!

