

UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Lecturer
Justin Yokota

Virtual Memory II

Page Table Details: Page Faults, Write Policy

- Page Table Details: Page Faults, Write Policy
- OS: Supervisor Mode, Exceptions
- Caches vs. Virtual Memory
- Translation Lookaside Buffer

Virtual Memory: Page Tables

Review

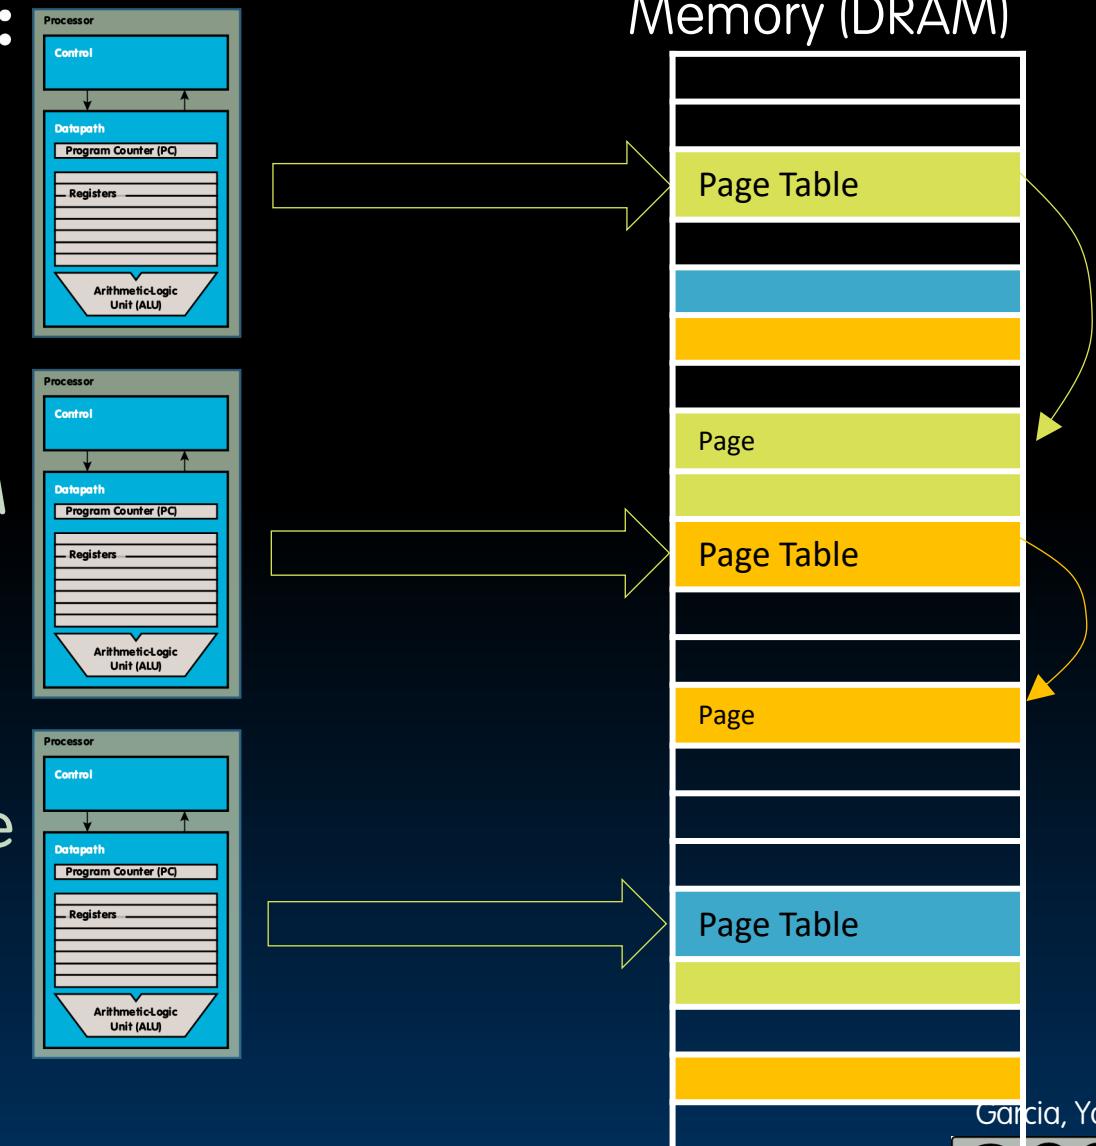
- Each process has a dedicated page table.
 - OS keeps track of which process is active.
- Isolation: Assign processes different pages in DRAM
 - Prevents (*protects*) a process from accessing other processes' data
 - Page tables managed by OS

We now have **two address spaces**:
virtual (per process) and physical (memory).



Page Tables Are Stored in Memory

- **lw/sw require two memory accesses:**
 - Read page table (stored in main memory) to translate to physical address.
 - Read physical page, also in main memory.
- **To minimize space and time penalty:**
 - Transfer blocks (not words) between DRAM and processor cache.
 - Have a hierarchical multi-level page table
 - Not in scope this semester, see CS162 for more
 - Use a cache for frequently used page table entries... (more later, TLB)



- Page table entries store status to indicate if the page is in memory (DRAM) or only on disk.
 - On each memory access, check the page table entry "*valid*" status bit.
- Valid → In DRAM
 - Read/write data in DRAM
- Not Valid → On disk
 - Triggers a *Page Fault*; OS intervenes to allocate the page into DRAM.
 - If out of memory, first evict a page from DRAM. ↘
 - Store evicted page to disk.
 - Read requested page from disk into DRAM.
 - Finally, read/write data in DRAM.

The *page replacement policy* (e.g., LRU/FIFO/random) is usually done in OS/software; this overheard << disk access time.

Memory's Write Policy?

- DRAM acts like a “cache” for disk.
 - Should writes always go directly to disk (write-through), or
 - Should writes only go to disk when page is evicted (write-back)?
- Answer: All virtual memory systems use write-back.
 - Disk accesses take too long!

Page Table Metadata: Status Bits

- **Write Protection Bit**

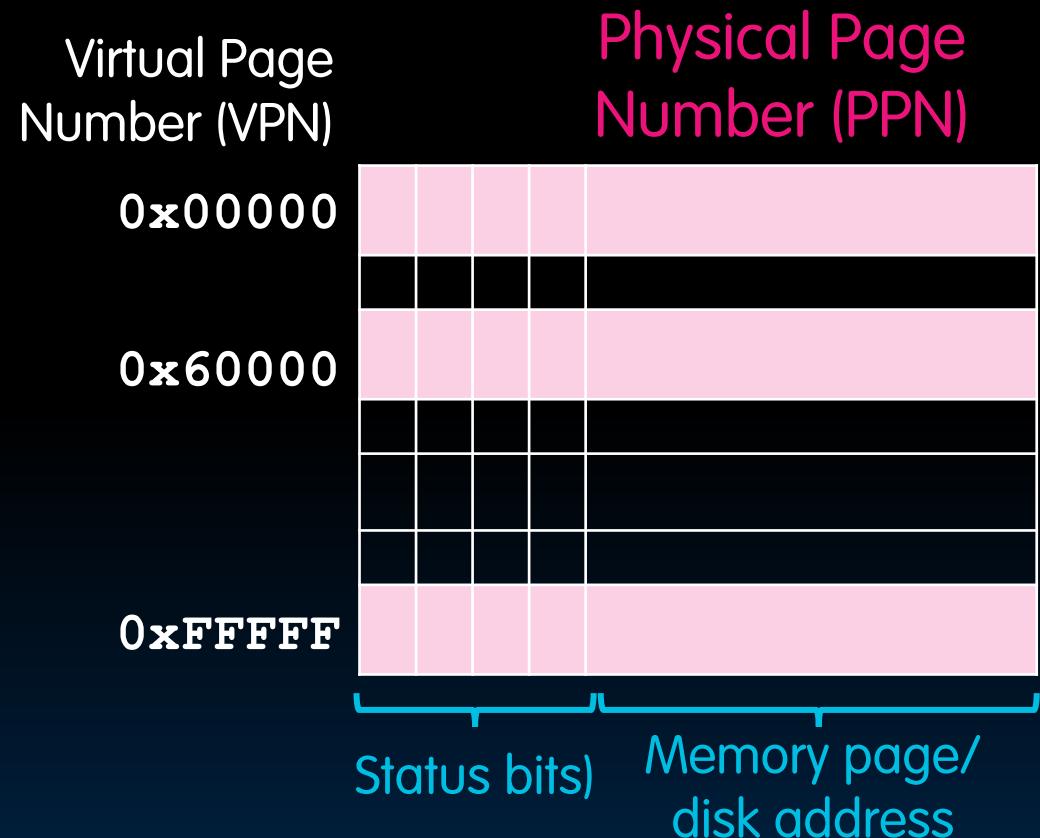
- On: If process writes to page, trigger exception

- **Valid Bit**

- On: Page is in RAM

- **Dirty Bit**

- On: Page on RAM is more up-to-date than page on disk



OS: Supervisor Mode, Exceptions

- Page Table Details: Page Faults, Write Policy
- OS: Supervisor Mode, Exceptions
- Caches vs. Virtual Memory
- Translation Lookaside Buffer

Supervisor Mode vs. User Mode

- If an application goes wrong (or rogue, e.g., malware), it could crash the entire machine!
- CPUs have a hardware supervisor mode (i.e., kernel mode).
 - Set by a status bit in a special register.
 - An OS process in supervisor mode helps enforce constraints to other processes, e.g., access to memory, devices, etc.
 - Supervisor mode is a bit like “superuser”...
 - Errors in supervisory mode are often catastrophic (blue “screen of death”, or “I just corrupted your disk”).
- By contrast, in user mode, a process can only access a subset of instructions and (physical) memory.
 - Can change *out* of supervisor mode using a special instruction (e.g. `sret`).
 - Cannot change *into* supervisor mode directly; instead, HW interrupt/exception.
 - The OS mostly runs in user mode! Supervisor mode is used sparingly.

Exceptions and Interrupts

▪ Exceptions

- Caused by an event *during* the execution of the current program.
- *Synchronous*; must be handled immediately.
- Examples:
 - Illegal instruction
 - Divide by zero
 - *Page fault*
 - Write protection violation

▪ Interrupts (more later)

- Caused by an event *external* to the current running program.
- *Asynchronous* to current program; does not need to be handled immediately (but should be soon).
- Examples:
 - Key press
 - Disk I/O

Traps Handle Exceptions/Interrupts

- The trap handler is code that services **interrupts/exceptions**.
asynchronous, synchronous,
external during (e.g. page fault)

1. Complete all instructions before the faulting instruction.
2. Flush all instructions after the faulting instruction.
 - Like pipeline hazard: convert to noops/"bubbles."
 - Also flush faulting instruction.
3. Transfer execution to trap handler (runs in **supervisor mode**).

- Optionally *return* to original program and re-execute instruction.



If the trap handler returns, then from the program's point of view it must look like nothing has happened!

, Yokota

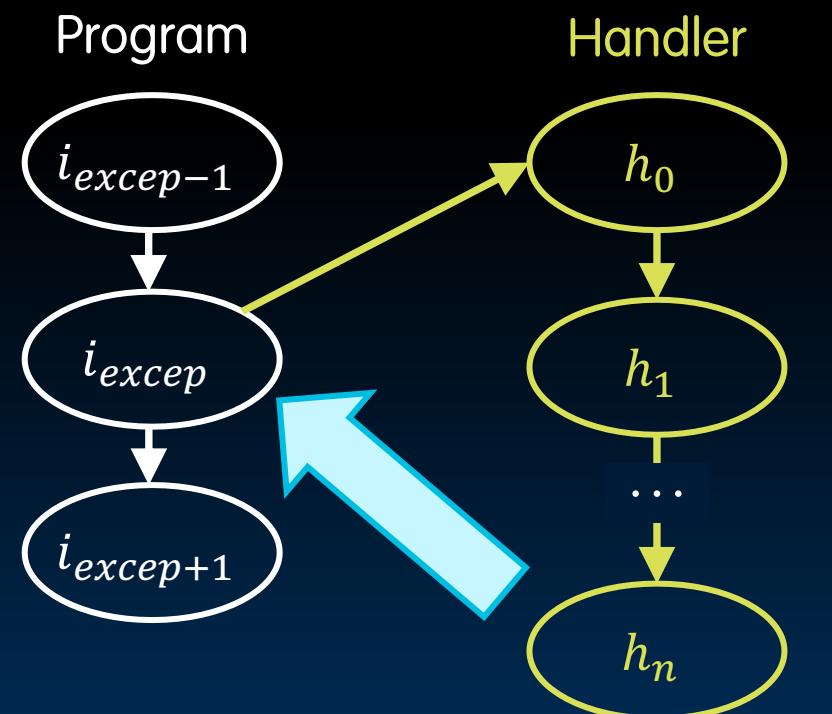
The Trap Handler

1. Save the state of the current program.
 - Save ALL of the registers!
2. Determine what caused the exception/interrupt.
3. Handle exception/interrupt, then do one of two things:



Continue execution of the program:

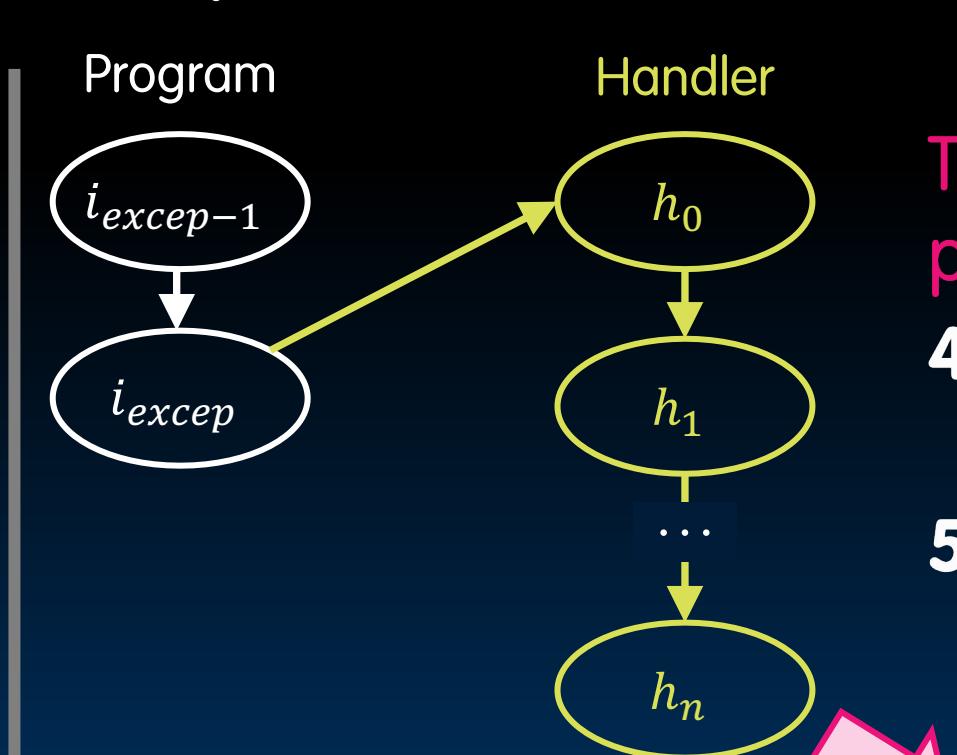
4. Restore program state.
5. Return control to the program.



The Trap Handler

1. **Save the state of the current program.**
 - Save ALL of the registers!
2. **Determine what caused the exception/interrupt.**
3. **Handle exception/interrupt, then do one of two things:**

Continue execution of the program:
4. Restore program state.
5. Return control to the program.



▫ Terminate the program:
4. Free the program resources, etc.
5. Schedule a new program.

Handling Context Switches

- **Recall the context switch:**
 - OS switches between processes (i.e., programs) by changing the internal state of the processor.
 - Allows a single processor to “simultaneously” run many programs.
- **At a high-level:**
 - The OS sets a timer. When it expires, perform a *hardware interrupt*.
 - Trap handler saves all register values, including:
 - Program Counter (PC)
 - *Page Table Register* (SPTBR in RV32I)
 - The memory *address* of the active process’s page table.
 - Trap handler then loads in the next process’s registers and returns to user mode.

Handling Page Faults

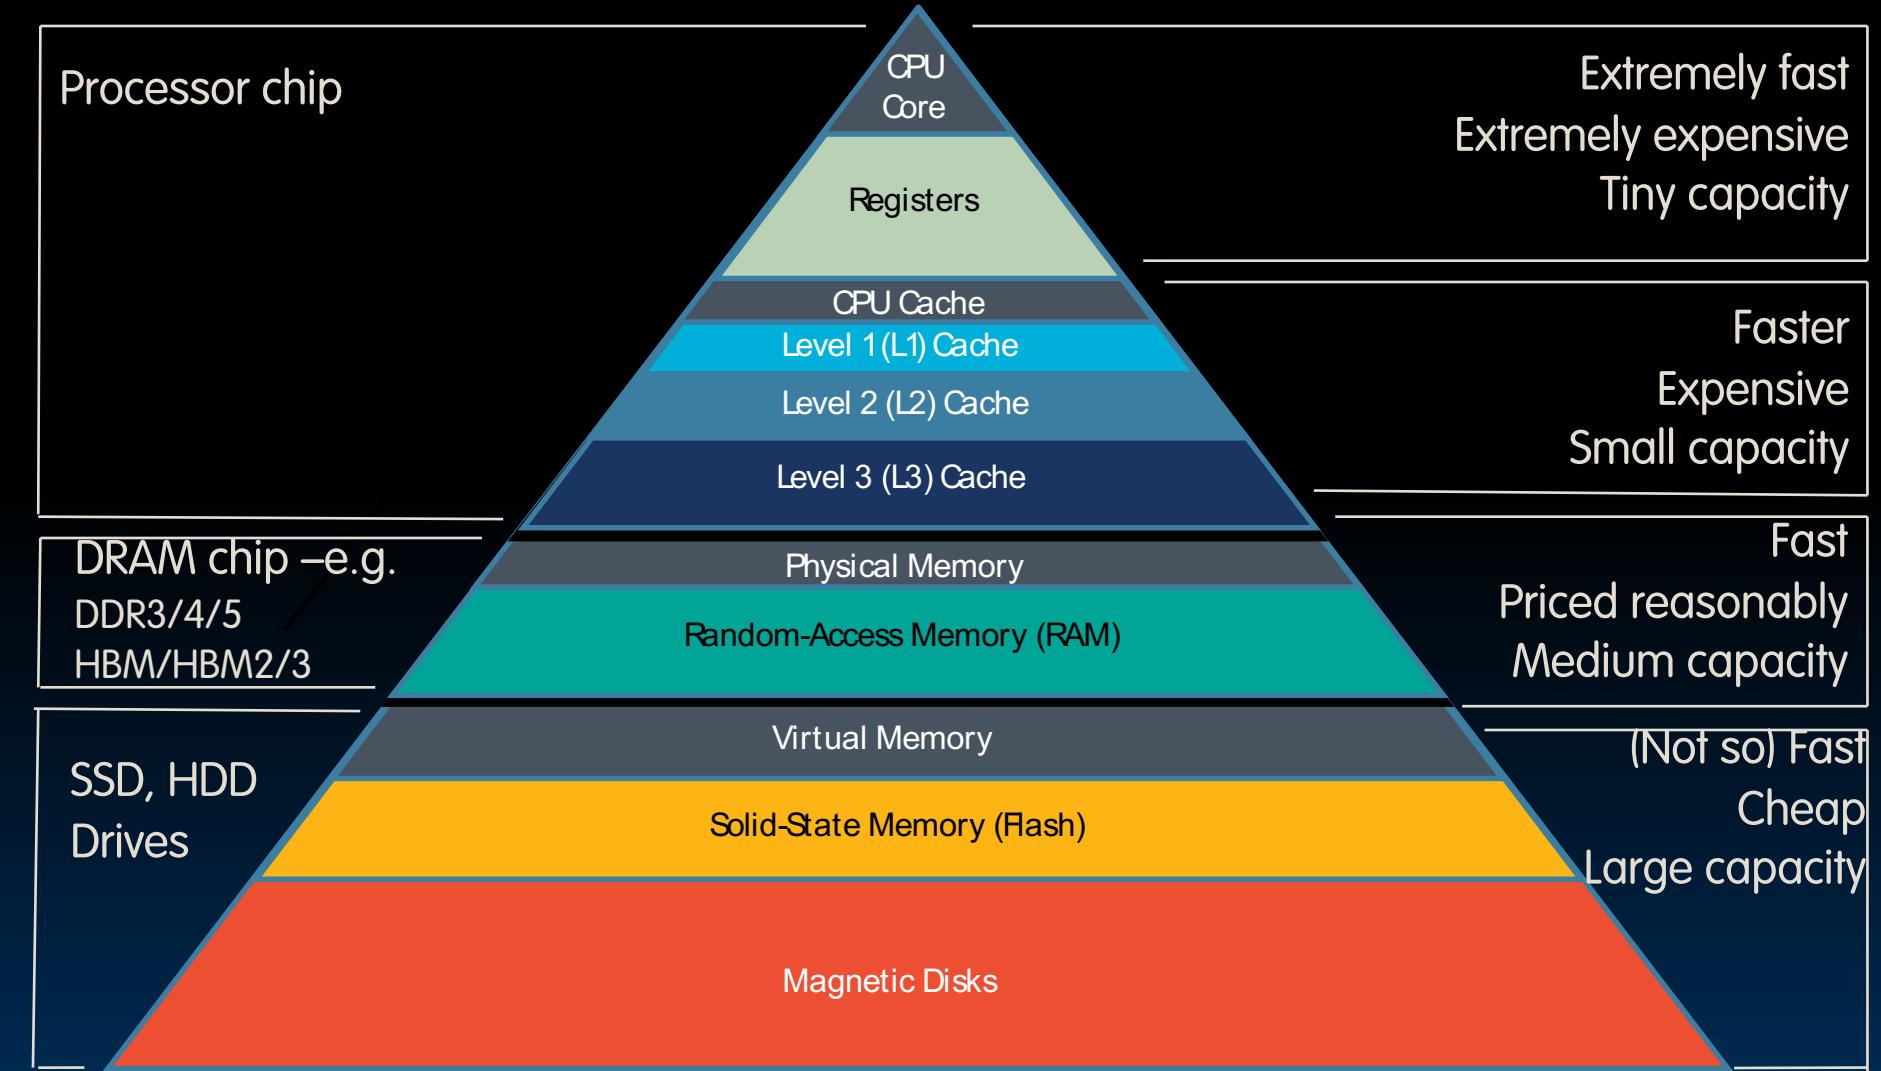
- **Recall page faults:**
 - An accessed page table entry has valid bit off → data is not in DRAM.
- **Page faults are handled by the trap handler.**
 - The *page fault exception handler* initiates transfers to/from disk and performs any page table updates.
 - (If pages needs to be swapped from disk, perform *context switch* so that another process can use the CPU in the meantime.)
 - (ideally need a “precise trap” [recoverable] so that resuming a process is easy.)
 - Following the page fault, *re-execute the instruction*.
- **Side note: Write protection violations also trigger exceptions.**

Caches vs. Virtual Memory

- Page Table Details: Page Faults, Write Policy
- OS: Supervisor Mode, Exceptions
- Caches vs. Virtual Memory
- Translation Lookaside Buffer

The Entire Modern Memory Hierarchy

Let's review the concepts of caches and memory.



Caches vs. Primary Memory

- **Blocks, pages, (bytes, words) are all units of memory.**
 - Caches: *blocks*
 - On modern systems, ~64B.
 - Memory: *pages*
 - On modern systems, ~4KiB.

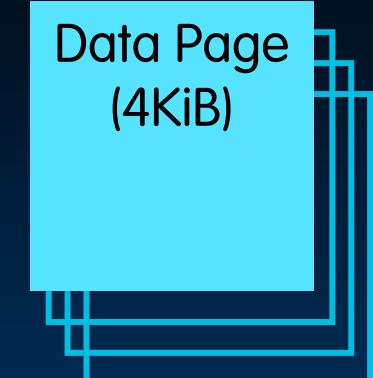
L1 Cache

V	Tag	Data
		Block (64B)
...

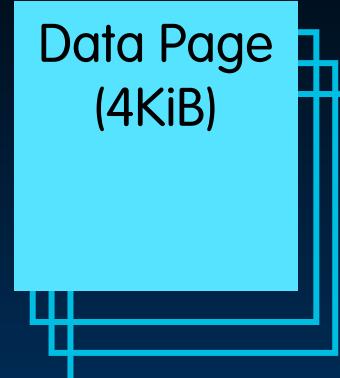
L2 Cache

V	Tag	Data
		Block (64B)
...

DRAM (primary memory)



Disk (secondary memory)



Caches vs. Page Tables

“Cache” Paradigm: Data at each level is a quick-access copy of data at a lower level in the memory hierarchy.

L1 Cache

V	Tag	Data
		😊
...

L2 Cache

V	Tag	Data
		😊
...

DRAM (primary memory)

Data Page



Similarly, DRAM data pages are “cached” disk pages.

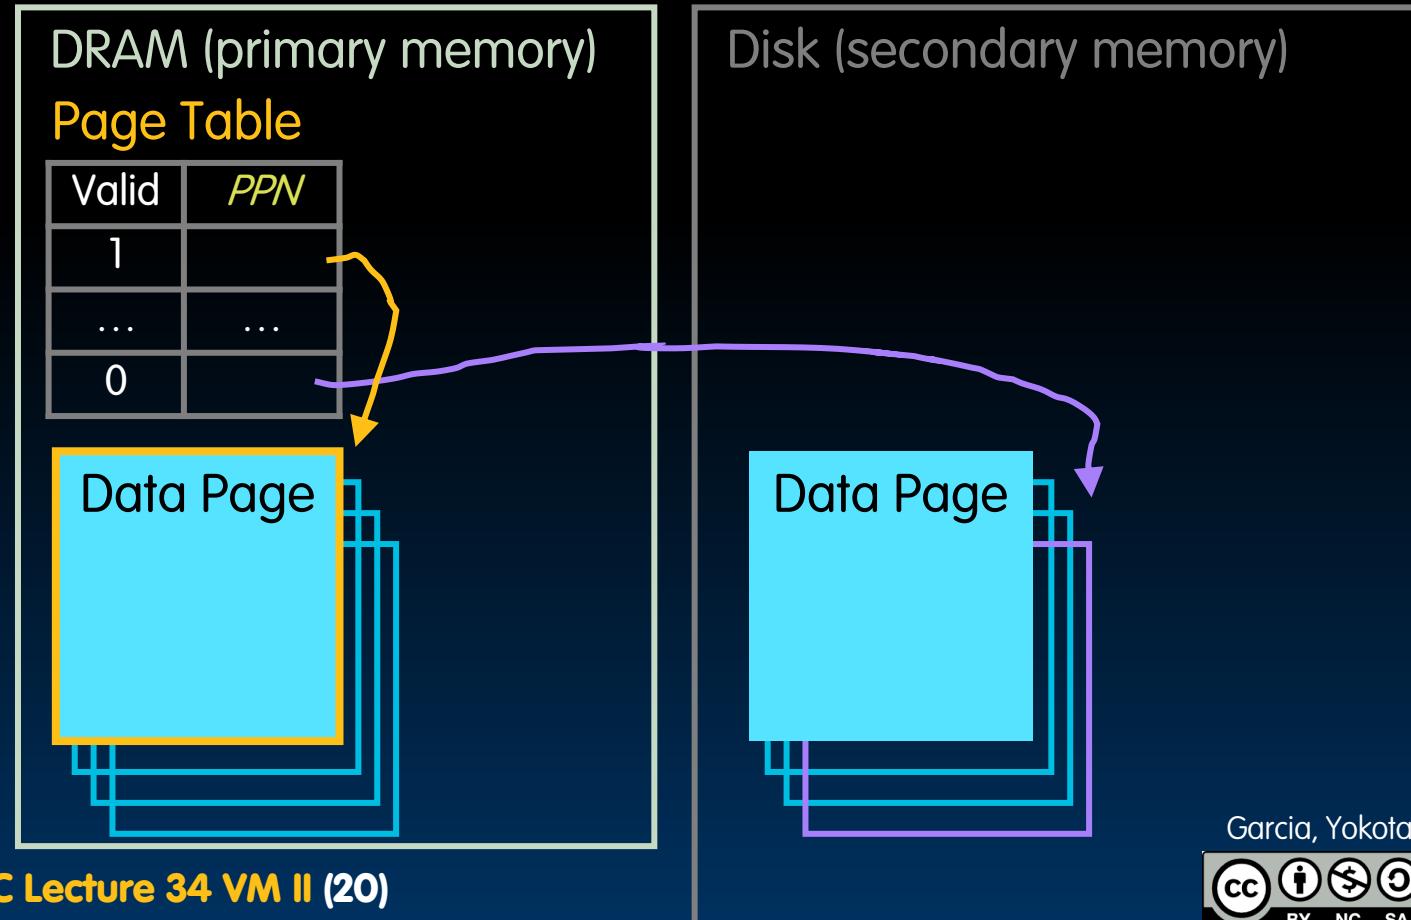
Data Page



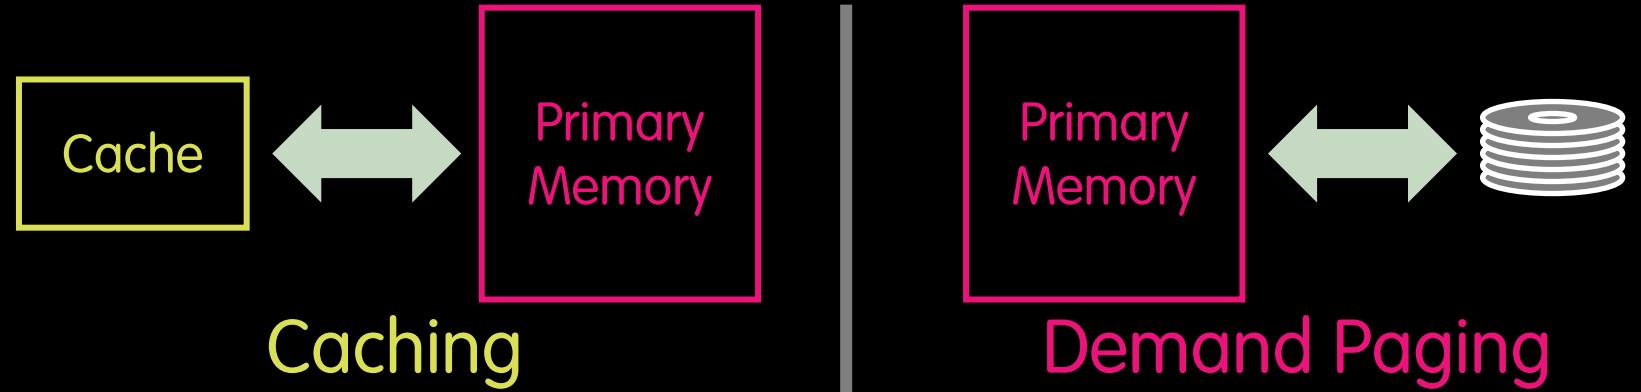
Caches vs. Page Tables

- A Page Table translates addresses.
 - Page tables store physical page numbers, *not data*.
- Page tables facilitate Demand Paging.
 - Cache data pages in memory.
 - Access disk pages only when needed by the process.
 - Page Table keeps track of page status/location.

(assuming single-level page tables)



Caching vs. Demand Paging



	Caching	Demand Paging
▫ Memory Unit	Block	Page
▪ Size	32B to 64B	4KiB to 8KiB
▪ Miss	Cache Miss	Page Fault
▫ Associativity	Direct-mapped, N-way Set associative, or fully associative	Fully associative (i.e., disk pages can be placed anywhere in DRAM)
▫ Replacement policy	Least Recently Used (LRU) or random	LRU (most common), or FIFO, or random
▫ Write policy	Write-through or write-back	Write-back

Translation Lookaside Buffer

- Page Table Details: Page Faults, Write Policy
- OS: Supervisor Mode, Exceptions
- Caches vs. Virtual Memory
- Translation Lookaside Buffer

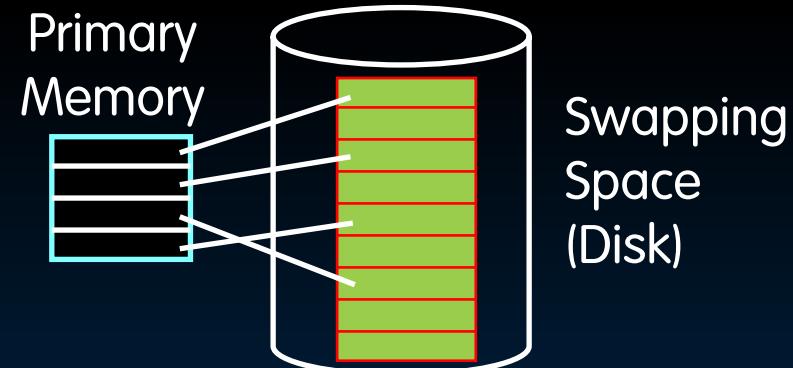
Modern Virtual Memory Systems

- Modern Virtual Memory Systems use address translation to provide the illusion of a large, private, and uniform storage.

1. Privacy means *Protection*:
 - Several users/processes, each with their own private address space.



2. Uniform storage means *Demand Paging*:
 - The ability to run programs larger than primary memory (DRAM).
 - Hides difference in machine configurations.



- Price: Address translation on each memory reference.

Page tables in memory significantly increase average memory access time!

Speeding Up Address Translation

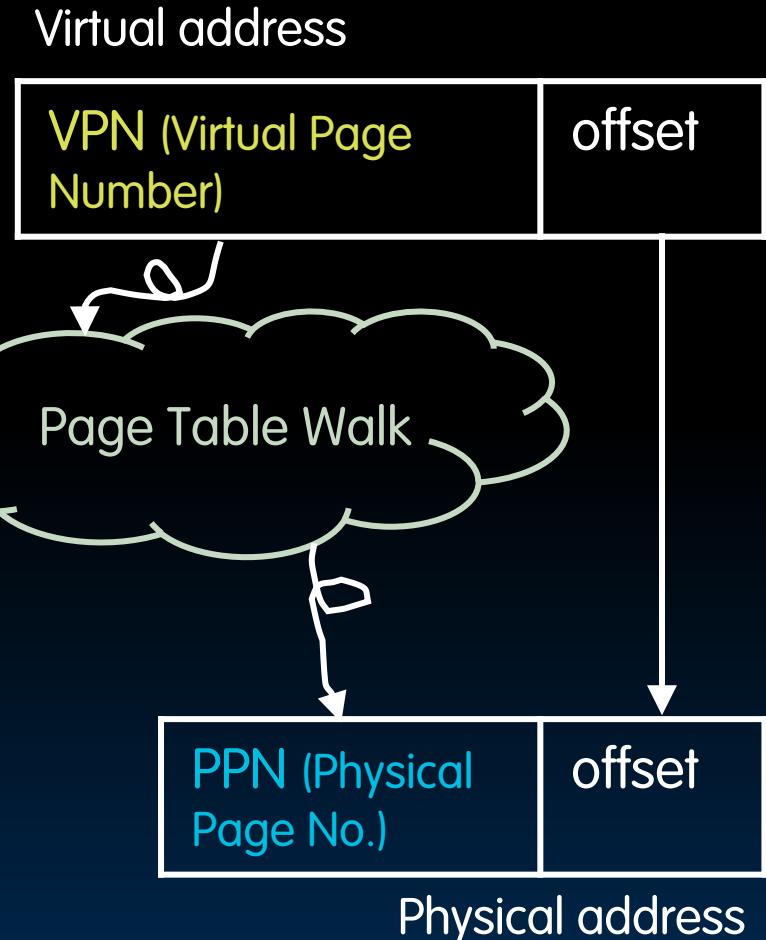
- Good Virtual Memory design should be fast (~1 clock cycle) and space efficient.

- Every instruction/data access needs address translation.

- But if page tables are in memory, then we must perform a page table walk per instruction/data access:

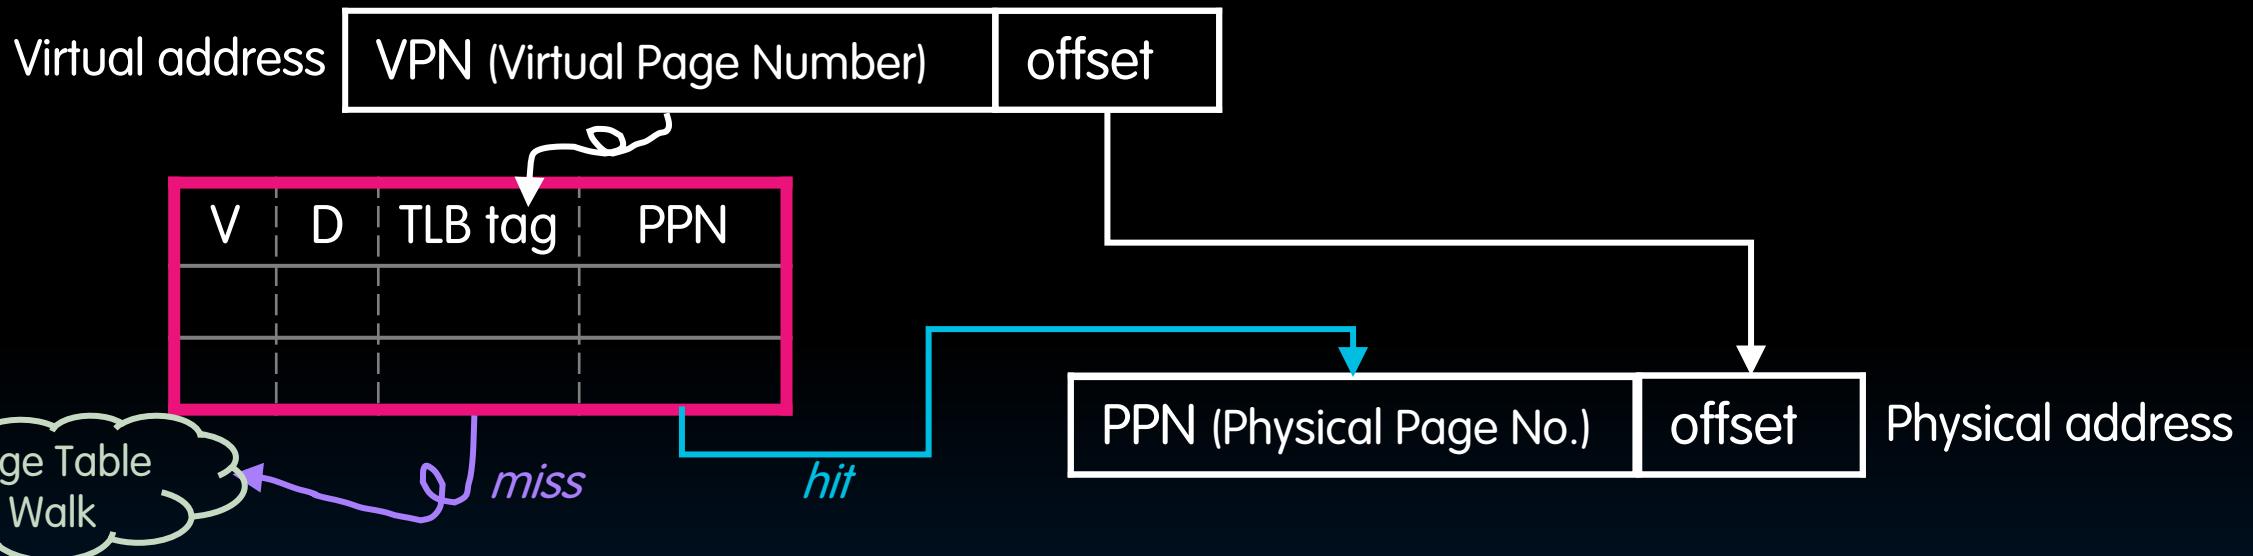
- Single-level page table: 2 memory accesses.
 - Two-level page table: 3 memory accesses.

- Solution: Cache some translations in the Translation Lookaside Buffer (TLB).



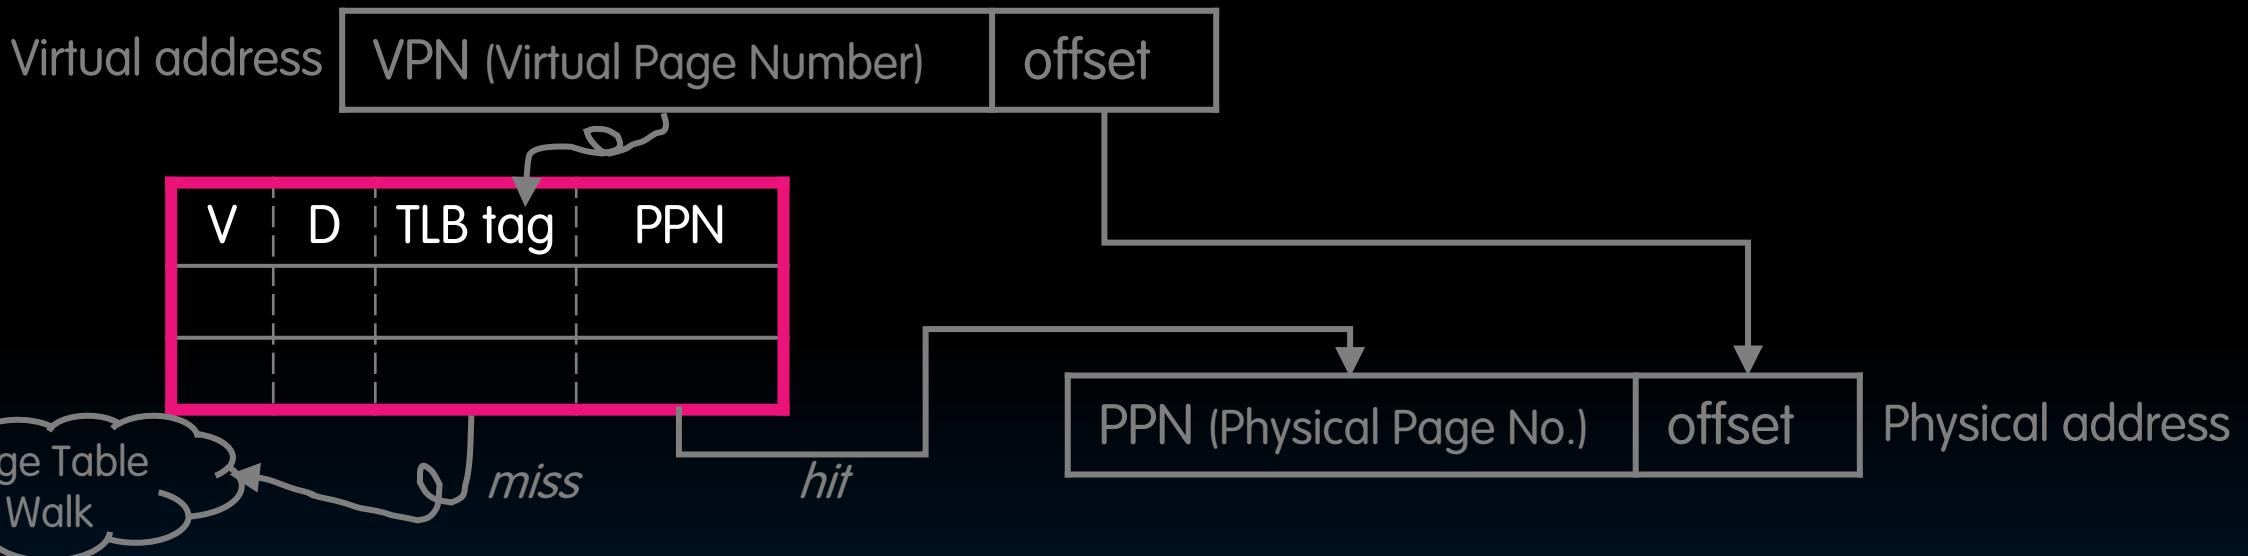
The TLB Is a Cache for Address Translations

- The Translation Lookaside Buffer (TLB) caches page table entries.
 - TLB *hit*: → Single-cycle translation ✓
 - TLB *miss*: → Page table walk to refill.



The TLB Is a Cache for Address Translations

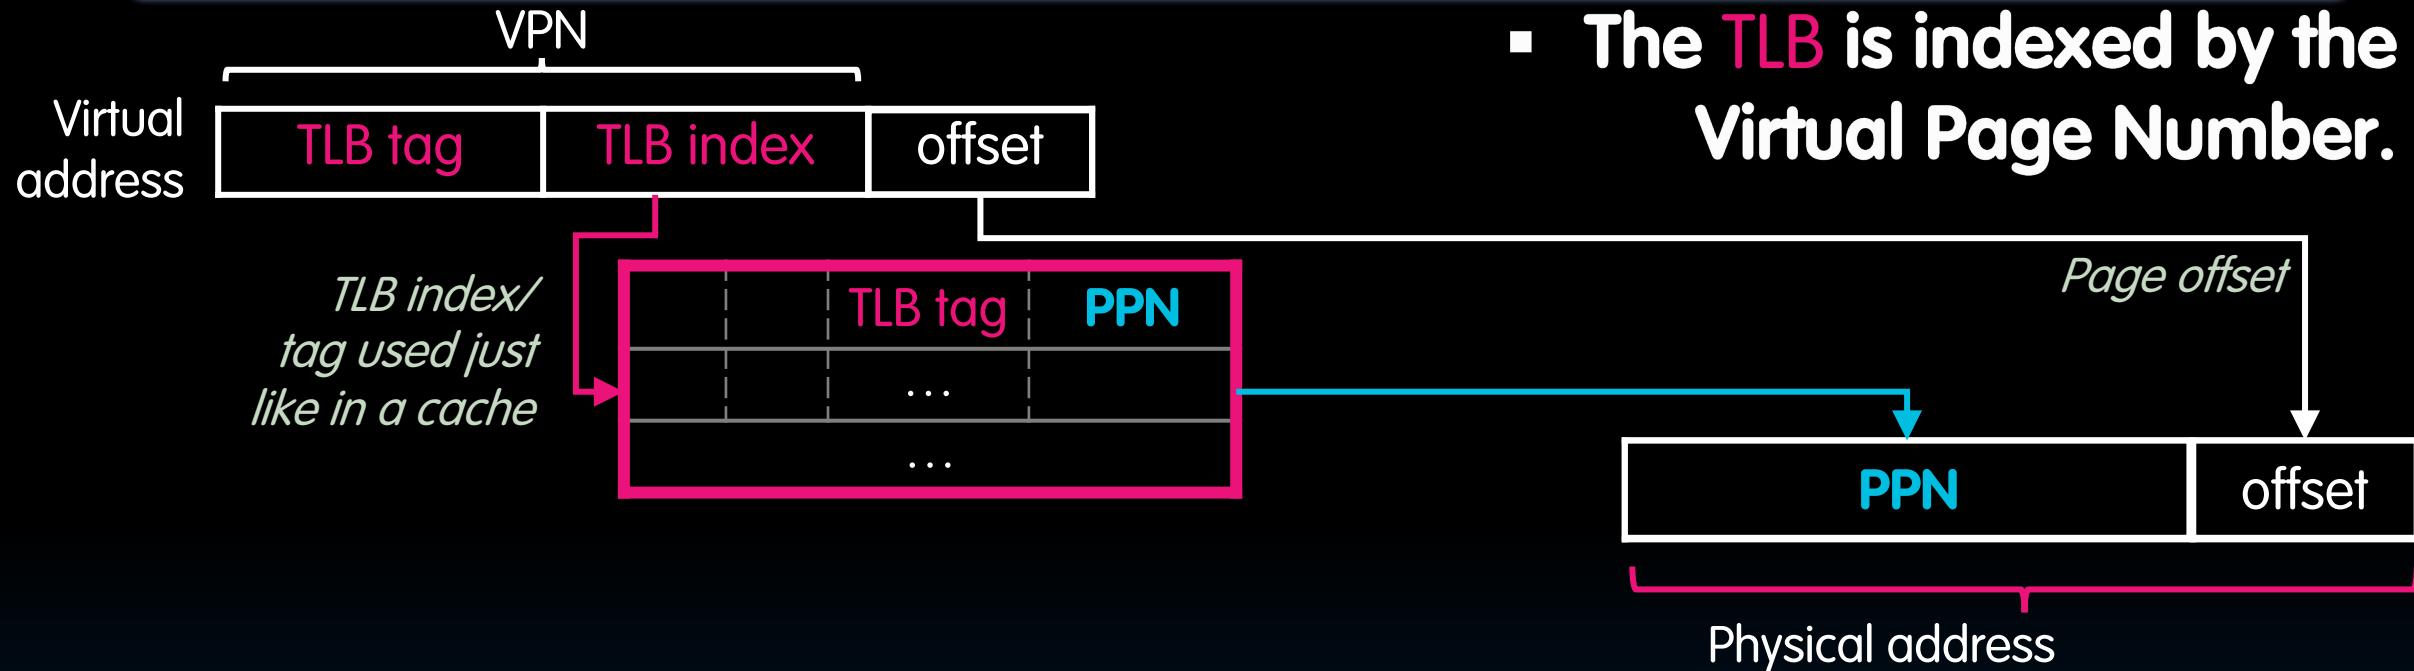
- The Translation Lookaside Buffer (TLB) **caches page table entries.**
 - TLB *hit*: → Single-cycle translation ✓
 - TLB *miss*: → Page table walk to refill.



- TLB Reach:** Size of largest virtual address space that can be simultaneously mapped by the TLB.
- TLB design:** 38-128 entries.
 - Typically *fully associative* (increase TLB reach by minimizing conflicting entries).
 - Random/FIFO replacement policy.

Tag, Index, and Offset

TIO for Virtual Addresses and Physical Addresses are *unrelated!*



Tag, Index, and Offset

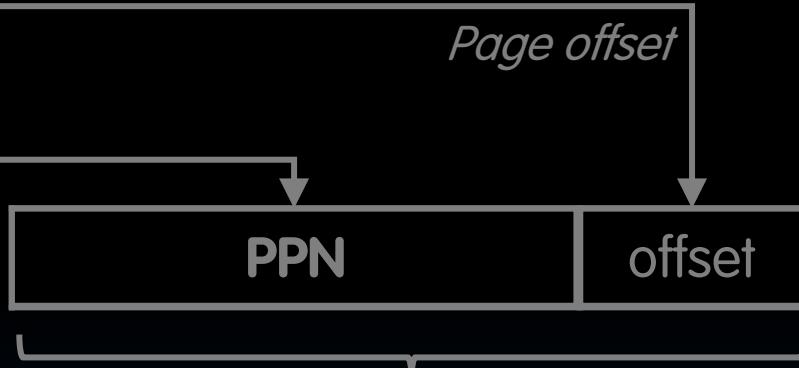
TIO for Virtual Addresses and Physical Addresses are *unrelated!*



*TLB index/
tag used just
like in a cache*



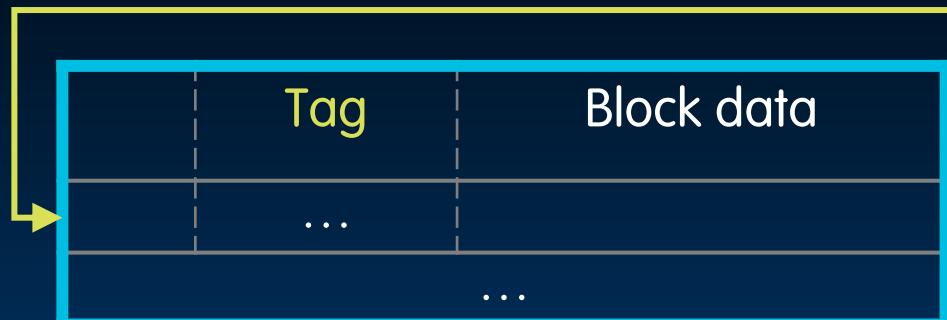
- The TLB is indexed by the Virtual Page Number.



Physical address (*split two ways*)

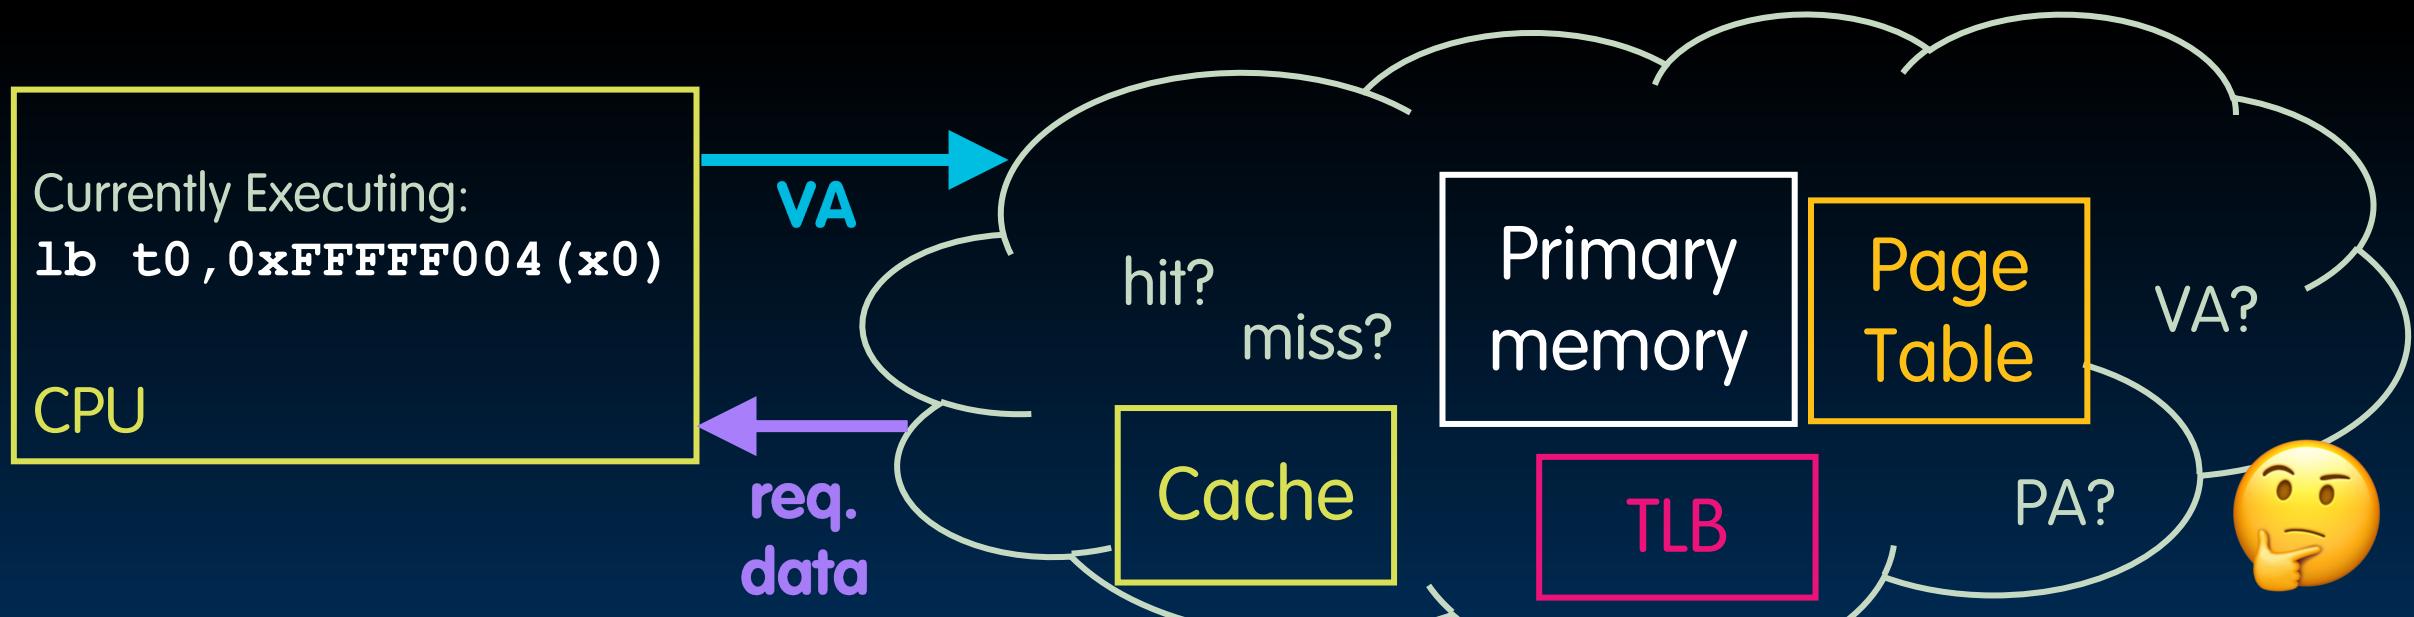


- The data cache is indexed by the Physical Address.



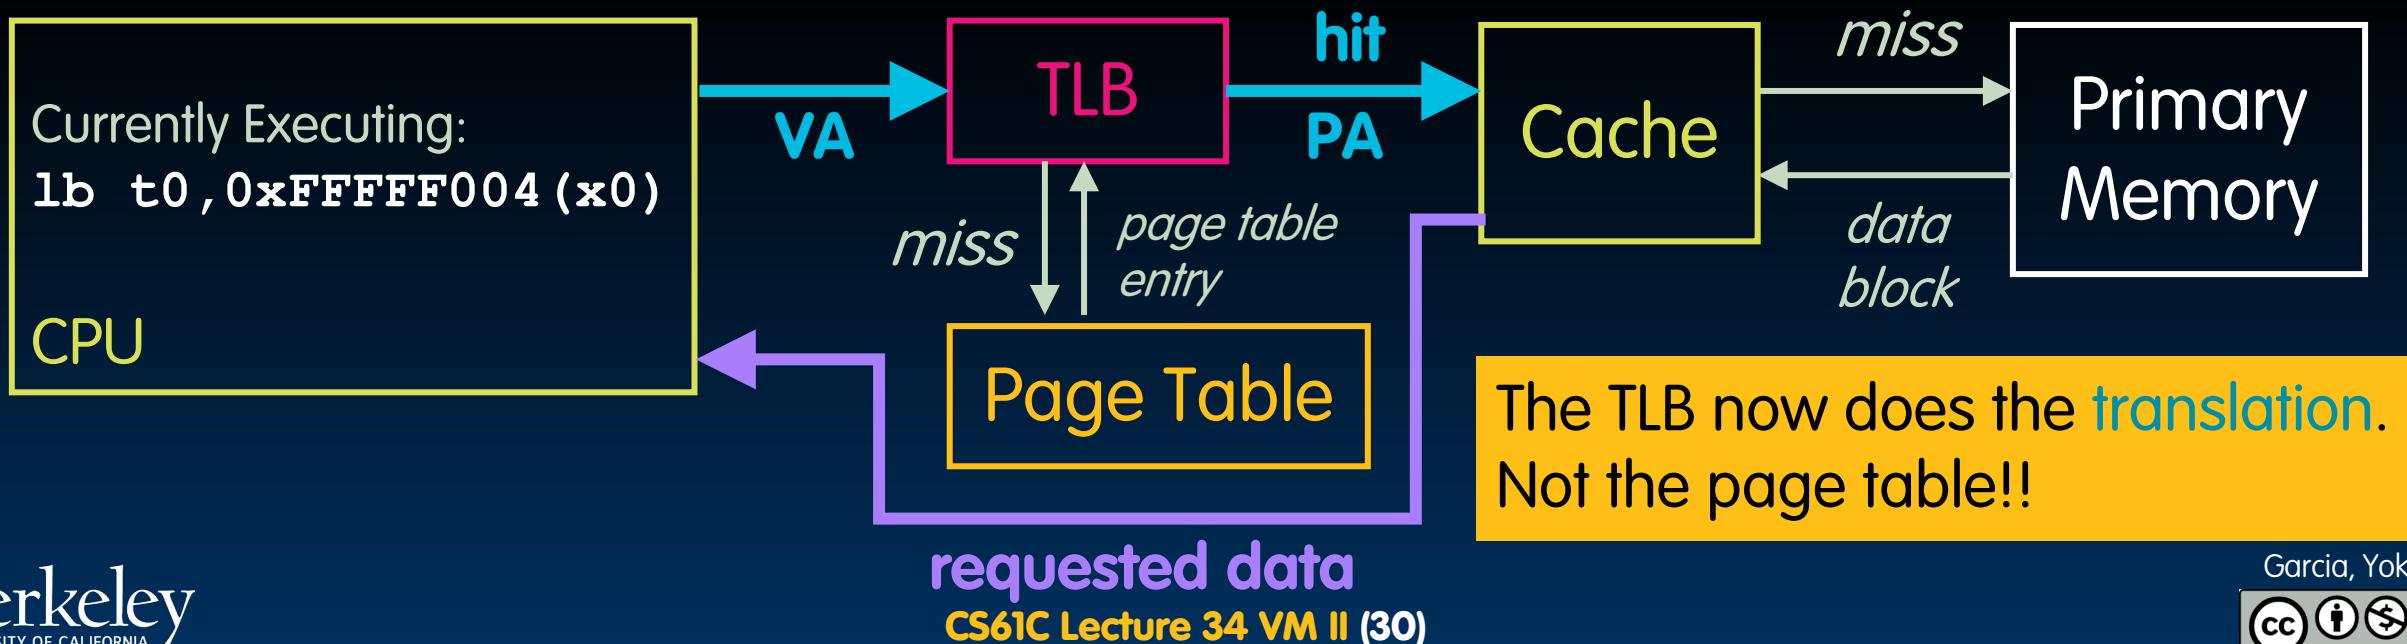
Memory Access: TLB, Cache, DRAM, Page Table?

1. Can a cache hold the requested data if the corresponding page is *not* in main memory?
2. On a memory reference, which block should we access first? When should we translate virtual addresses?



Memory Access: TLB, Cache, DRAM, Page Table

1. Can a cache hold the requested data if the corresponding page is *not* in main memory? No!
2. On a memory reference, which block should we access first? When should we translate virtual addresses?
 - We will assume *Physically Indexed, Physically Tagged* caches (other designs exist).
 - This means TLB first, then cache.



- **Virtual Memory is a wonderful abstraction!**
 - Run multiple processes, each with independent **protected** 4GiB memory
 - Have physical memory **smaller than or bigger than** virtual memory
 - Page tables live in memory, time and space expensive
 - Valid (in memory vs on disk) and **Dirty** status bits (write back!)
- **Trap Handler runs in Supervisor Mode**
 - Handles **exceptions** (sync, e.g., page faults) or **interrupts** (async)
 - Swaps pages out on page fault with context switch so no lost time
- **Caches are caches for memory, Memory is cache for disk**
- **TLB is a cache of the page table**
 - Most recent page table entries stored
- **Cache is physically addressed (uses PA not VA)**
- **Want to make your computer faster?**
 - Buy more memory! (can run more programs “resident”, avoid going to Pluto)