

UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley  
Lecturer  
Justin Yokota

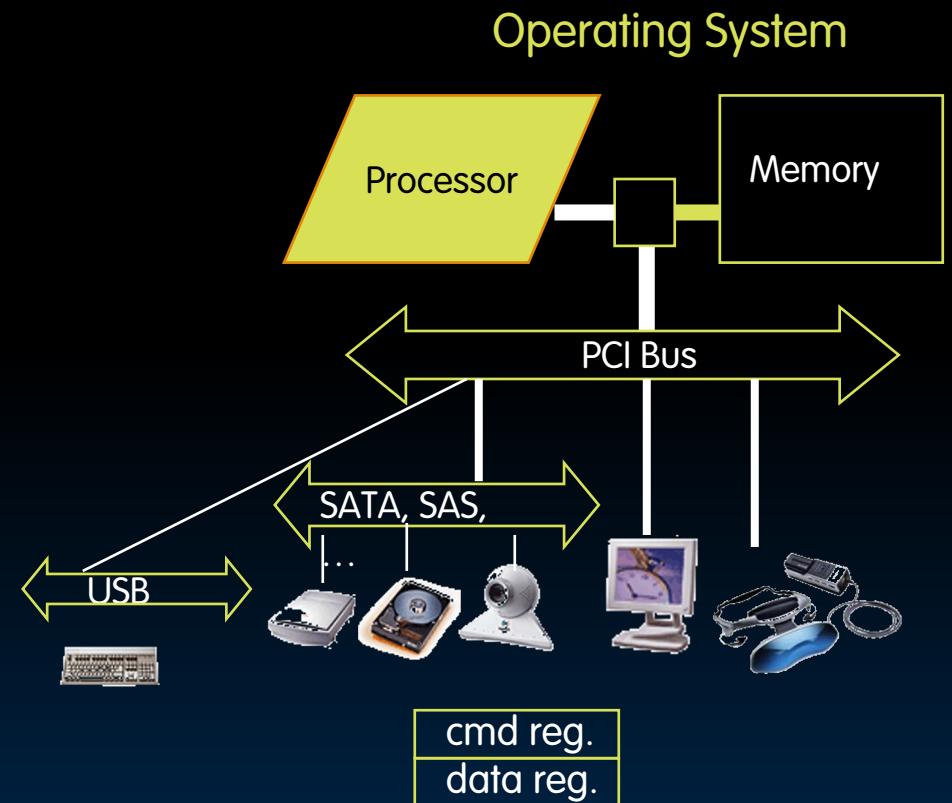
## Input/Output (I/O)

# I/O Devices

- I/O Devices
- I/O Polling
- I/O Interrupts
- Direct Memory Access (DMA)
- Networking

# How to Interact with Devices?

- Assume a program running on a CPU. How does it interact with the outside world?
- Need I/O interface for keyboards, network, mouse, display, etc.
  - Connect to many types of devices
  - Control these devices, respond to them, and transfer data
  - Present them to user programs so they are useful



- **What must the processor do for I/O?**

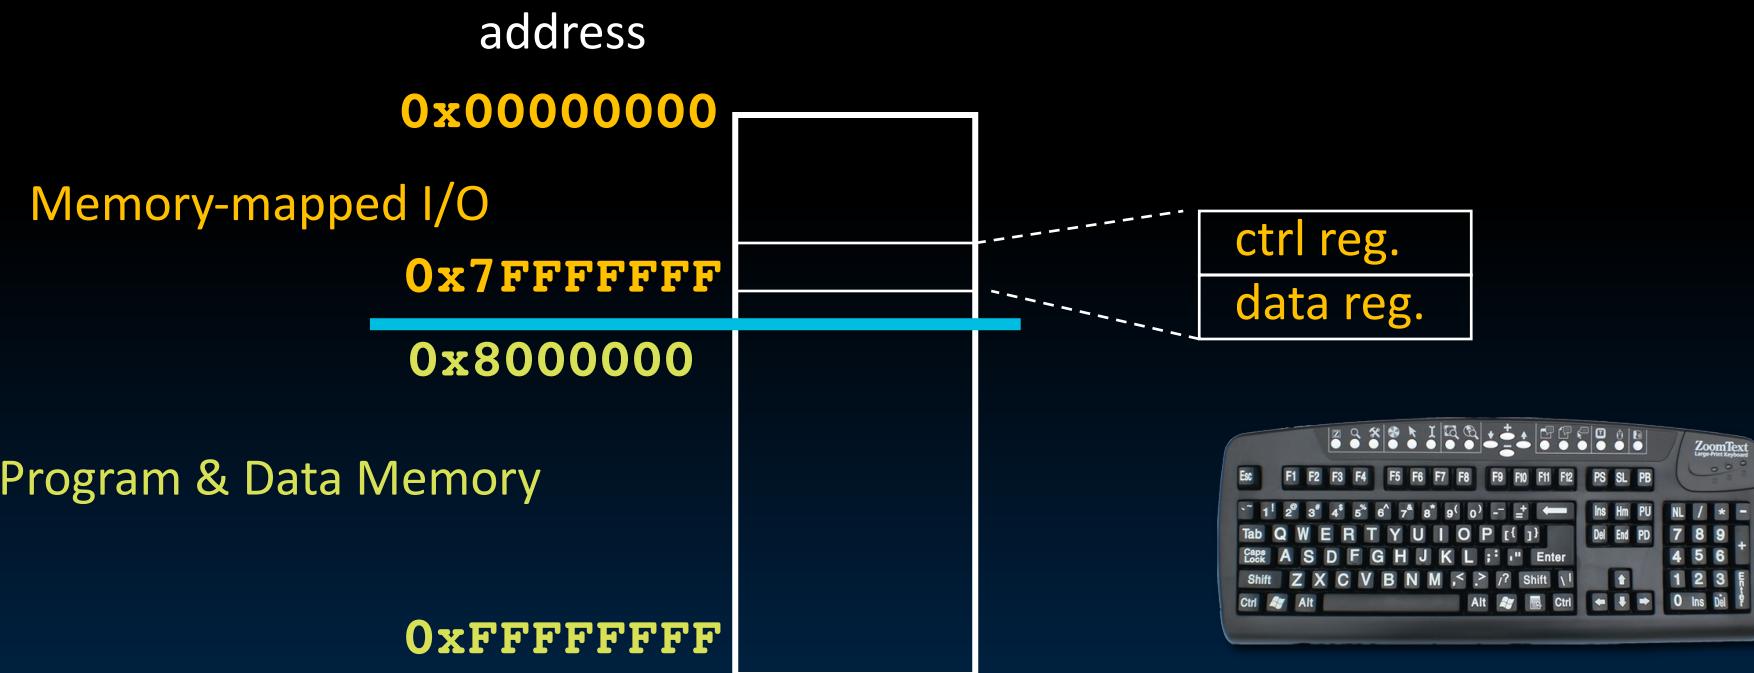
- Input: Read a sequence of bytes
  - Output: Write a sequence of bytes

- **Interface options**

- a) Special input/output instructions & hardware
  - b) Memory mapped I/O
    - Portion of address space dedicated to I/O
    - I/O device registers there (no memory)
    - Use normal load/store instructions, e.g. **lw / sw**
    - Very common, used by RISC-V

# Memory Mapped I/O

- Certain addresses are not 'regular memory'
- Instead, they correspond to registers in I/O devices



# Processor-I/O Speed Mismatch

- **1 GHz microprocessor I/O throughput:**
  - 4 GiB/s (lw/sw)
- **Typical I/O data rates:**
  - 10 B/s (keyboard)
  - 3 MiB/s (Bluetooth 3.0)
  - 0.06-1.25 GiB/s (USB 2/3.1)
  - 7-250 MiB/s (Wifi, depends on standard)
  - 125 MiB/s (G-bit Ethernet)
  - 480 MiB/s (SATA3 HDD)
  - 560 MiB/s (cutting edge SSD)
  - 5 GiB/s (Thunderbolt 3)
  - 32 GiB/s (High-end DDR4 DRAM)
  - 64 GiB/s (HBM2 DRAM)
  - These are peak rates – actual throughput is lower
- **Common I/O devices neither deliver nor accept data matching processor speed**

# I/O Polling

- I/O Devices
- I/O Polling
- I/O Interrupts
- Direct Memory Access (DMA)
- Networking



# Polling: Processor Checks Status, Then Acts

- Device registers generally serve two functions:
  - *Control Register*, says it's OK to read/write (I/O ready) [think of a flagman on a road]
  - *Data Register*, contains data
- Processor reads from Control Register in loop
  - Waiting for device to set Ready bit in Control reg ( $0 \rightarrow 1$ )
  - Indicates "data available" or "ready to accept data"
- Processor then loads from (input) or writes to (output) data register
  - I/O device resets control register bit ( $1 \rightarrow 0$ )
- Procedure called "Polling"

# I/O Example (Polling)

- Input: Read from keyboard into a0

```
lui t0 0x7fffff #7fffff000 (io addr)
```

**Waitloop:** lw t1 0(t0) #read control Memory map

```
andi t1 t1 0x1 #ready bit
```

```
beq t1 zero Waitloop
```

```
lw a0 4(t0) #data
```

- Output: Write to display from a1

```
lui t0 0x7fffff #7fffff000
```

**Waitloop:** lw t1 8(t0) #write control

```
andi t1 t1 0x1 #ready bit
```

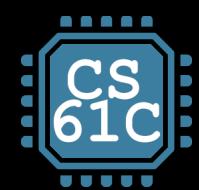
```
beq t1 zero Waitloop
```

```
sw a1 12(t0) #data
```

- “Ready” bit is from processor’s point of view!

# Cost of Polling?

- **Assume for a processor with**
  - 1 GHz clock rate
  - Taking 400 clock cycles for a polling operation
  - Call polling routine
    - Check device (e.g., keyboard or WiFi input available)
    - Return
    - What's the percentage of processor time spent polling?
- **Example:**
  - Mouse
  - Poll 30 times per second
    - Set by requirement not to miss any mouse motion (which would lead to choppy motion of the cursor on the screen)



# % Processor Time to Poll Mouse

- **Mouse Polling [clocks/sec]**
  - $= 30 \text{ [polls/s]} * 400 \text{ [clocks/poll]} = 12K \text{ [clocks/s]}$
- **% Processor for polling:**
  - $12*10^3 \text{ [clocks/s]} / 1*10^9 \text{ [clocks/s]} = 0.0012\%$   
→ Polling mouse little impact on processor...
- **(Except that we need to know we should be polling...)**
  - ...remember the timer for context switches? Could use that

# % Processor Time to Poll Hard Disk

- Frequency of Polling Disk (rate at which chunks come off disk)
  - $= 16 \text{ [MB/s]} / 16 \text{ [B/poll]} = 1M \text{ [polls/s]}$
- Disk Polling, Clocks/sec =  
 $1M \text{ [polls/s]} * 400 \text{ [clocks/poll]} = 400M \text{ [clocks/s]}$
- % Processor for polling:
  - $400*10^6 \text{ [clocks/s]} / 1*10^9 \text{ [clocks/s]} = 40\%$   
→ Unacceptable
    - (Polling is only part of the problem – accessing in small chunks is inefficient, too)

# I/O Interrupts

- I/O Devices
- I/O Polling
- I/O Interrupts
- Direct Memory Access (DMA)
- Networking

# Alternatives to Polling: Interrupts

- **Polling wastes processor resources**
  - Akin to waiting at the door for guests to show up
  - What about a bell?
- **Computer lingo for bell:**
  - Interrupt
  - Occurs when I/O is ready or needs attention
    - **Interrupt current program**
    - Transfer control to the trap handler in the operating system
- **Interrupts:**
  - No I/O activity: Nothing to do
  - Lots of I/O: Expensive – thrashing caches, VM, saving/restoring state

# Polling, Interrupts and DMA

- **Low data rate (e.g. mouse, keyboard)**
  - Keyboard uses interrupts or polling (depending on the HW interface)
    - Overhead of interrupts ends up being low
  - Mouse input is typically polled
- **High data rate (e.g. network, disk)**
  - Start with interrupts...
    - If there is no data, you don't do anything!
  - Once data starts coming... Switch to Direct Memory Access (DMA)

# Aside: Programmed I/O

- “Programmed I/O”:
  - Standard for ATA hard-disk drives
  - CPU execs lw/sw instructions for all data movement to/from devices
  - CPU spends time doing two things:
    - Getting data from device to main memory
    - Using data to compute
- Not ideal because ...
  1. CPU has to execute all transfers, could be doing other work
  2. Device speeds don’t align well with CPU speeds
  3. Energy cost of using beefy general-purpose CPU where simpler hardware would suffice
- Until now CPU has sole control of main memory
- 5% of CPU cycles on Google Servers spent in `memcpy()` and `memmove()` library routines!\*

\*Kanev et al., “Profiling a warehouse-scale computer,” ICSA 2015, (June 2015), Portland, OR.

# Direct Memory Access (DMA)

- I/O Devices
- I/O Polling
- I/O Interrupts
- Direct Memory Access (DMA)
- Networking

# Direct Memory Access (DMA)

- Allows I/O devices to directly read/write main memory
- New hardware: The **DMA Engine**
- DMA engine contains registers written by CPU:
  - Memory address to place data
  - # of bytes
  - I/O device #, direction of transfer
  - unit of transfer, amount to transfer per burst

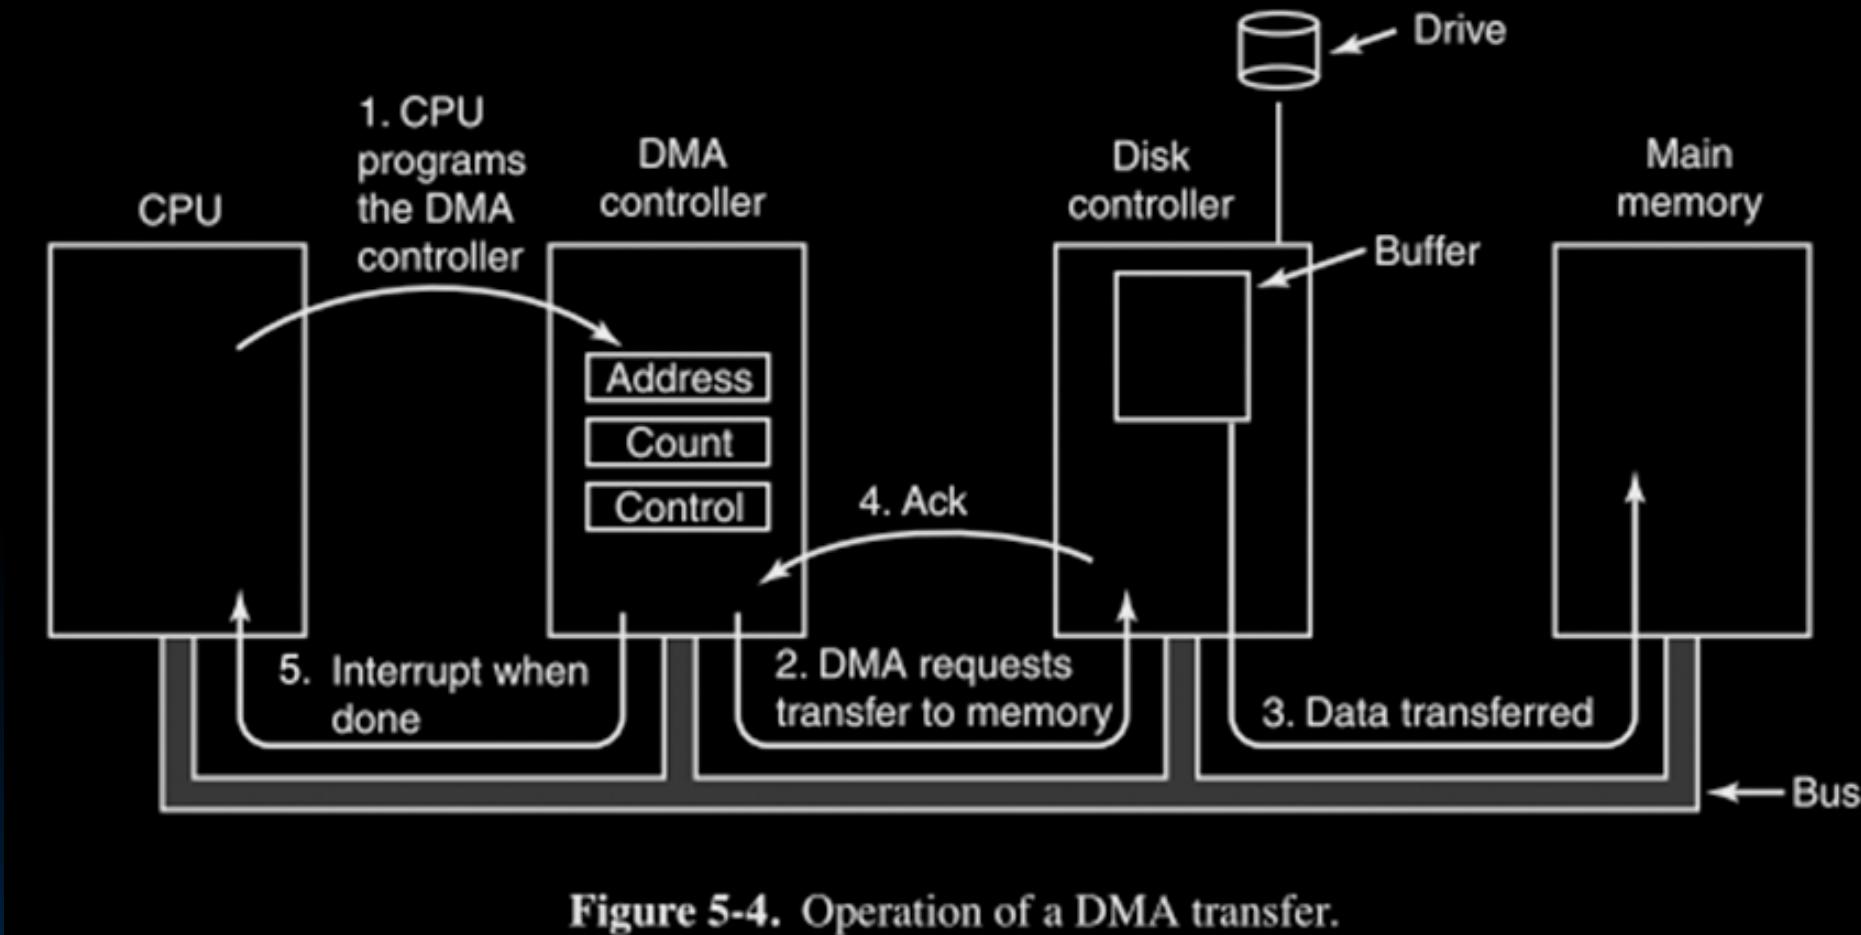


Figure 5-4. Operation of a DMA transfer.

From Section 5.1.4 Direct Memory Access in *Modern Operating Systems* by Andrew S. Tanenbaum, Herbert Bos, 2014

# DMA: Incoming Data

1. Receive interrupt from device
2. CPU takes interrupt, initiates transfer
  - Instructs DMA engine/device to place data @ certain address
3. Device/DMA engine handle the transfer
  - CPU is free to execute other things
4. Upon completion, Device/DMA engine interrupt the CPU again

# DMA: Outgoing Data

1. CPU decides to initiate transfer, confirms that external device is ready
2. CPU begins transfer
  - Instructs DMA engine/device that data is available @ certain address
3. Device/DMA engine handle the transfer
  - CPU is free to execute other things
4. Device/DMA engine interrupt the CPU again to signal completion

# Networking

- I/O Devices
- I/O Polling
- I/O Interrupts
- Direct Memory Access (DMA)
- Networking

- Originally sharing I/O devices between computers
  - E.g., printers
- Then communicating between computers
  - E.g., file transfer protocol
- Then communicating between people
  - E.g., e-mail
- Then communicating between networks of computers
  - E.g., file sharing, www, ...

# The Internet (1962)

## ▪ History

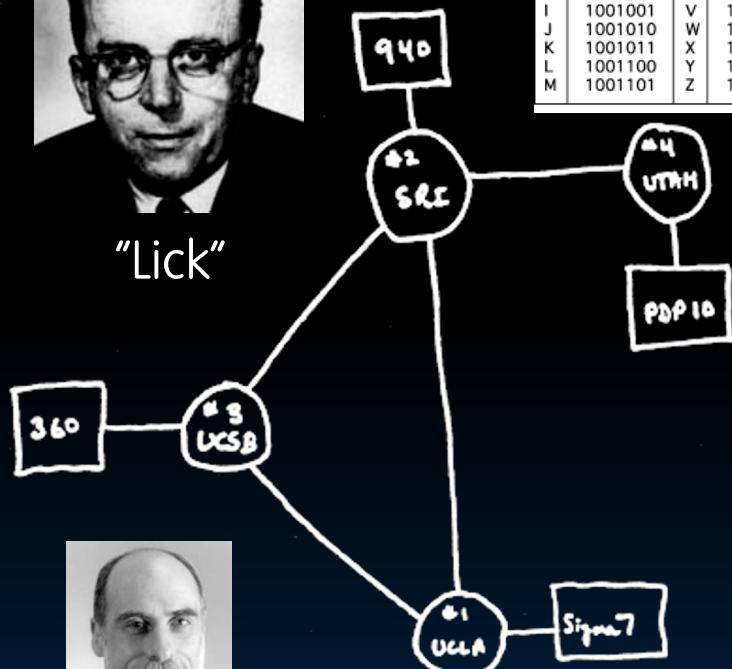
- 1963: JCR Licklider, while at DoD's ARPA, writes a memo describing desire to connect the computers at various research universities: Stanford, Berkeley, UCLA, ...
- 1969 : ARPA deploys 4 "nodes" @ UCLA, SRI, Utah, & UCSB
- 1973 Robert Kahn & Vint Cerf invent TCP, now part of the Internet Protocol Suite

## ▪ Internet growth rates

- Exponential since start!



"Lick"



Vint Cerf

ASCII Alphabet			
A	1000001	N	1001110
B	1000010	O	1001111
C	1000011	P	1010000
D	1000100	Q	1010001
E	1000101	R	1010010
F	1000110	S	1010011
G	1000111	T	1010100
H	1001000	U	1010101
I	1001001	V	1010110
J	1001010	W	1010111
K	1001011	X	1011000
L	1001100	Y	1011001
M	1001101	Z	1011010

# The World Wide Web (1989)

- “System of interlinked hypertext documents on the Internet”

- **History**

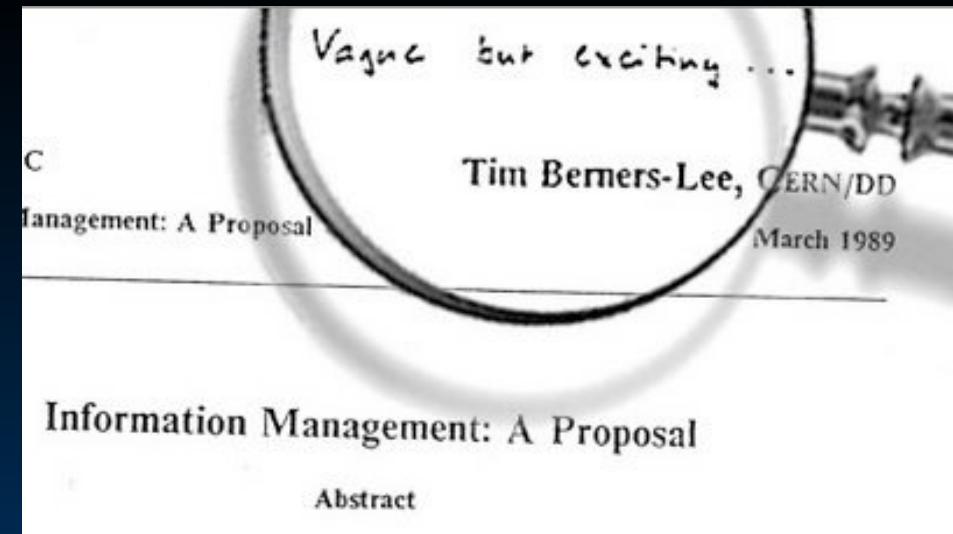
- 1945: Vannevar Bush describes hypertext system called “memex” in article
  - 1989: Sir Tim Berners-Lee proposed and implemented the first successful communication between a Hypertext Transfer Protocol (HTTP) client and server using the internet.
  - ~2000 Dot-com entrepreneurs rushed in, 2001 bubble burst

- **Today : Access anywhere!**



World's First web server in 1990

Tim Berners-Lee



Information Management: A Proposal

Abstract

Garcia, Yokota



## ■ SW Send steps

- 1: Application copies data to OS buffer
- 2: OS calculates checksum, starts timer
- 3: OS sends data to network interface HW and says start

## ■ SW Receive steps

- 3: OS copies data from network interface HW to OS buffer
- 2: OS calculates checksum, if OK, send ACK; if not, delete message (sender resends when timer expires)
- 1: If OK, OS copies data to user address space, & signals application to continue



# What do we need?

- Traditionally, a Network Interface Card (NIC)



- Wired or wireless
- Transfers data by using programmed I/O (old) or DMA (new)

🌐 When poll is active, respond at [pollev.com/ddg](https://pollev.com/ddg)

SMS Text **DDG** to 22333 once to join

## L35 Three I/O questions

- Faster CPU results in faster I/O
- Low-level I/O is actually quite simple, since it's just reading/writing bytes
- Packets can take different paths and arrive out of order

FFF  
FFT  
FTF  
FTT  
TFF  
TFT  
TTF  
TTT

Powered by  Poll Everywhere

Start the presentation to see live content. For screen share software, share the entire screen. Get help at [pollev.com/app](https://pollev.com/app)

# "And in Conclusion..."

- **We have figured out how computers work!**
  - And figured out how the OS works and how to interact with it
- **We have built a virtual memory system**
  - And have developed understanding of physical memory, storage devices
- **And we can attach peripherals for I/O!**